Fine-tuning an LLM

What does pre-training mean?

What does fine-tuning mean?

How many parameters does BERT have?

Is BERT much smaller than GPT?

How was the BERT model pre-trained?

How does MLM pre-training objective work?

How does NSP pre-training objective work?

Pre-training

MLM on unlabelled data

word2vec
GloVe
skip-thought
InferSent
ELMo
ULMFiT
GPT
BERT

Fine-tuning

Cross-entropy on task labels

classification
sequence labeling
Q&A
....

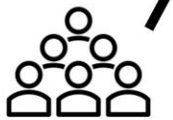

Pre-training

is like a child learning to
read and write his/her mother tongue.

Fine Tuning

is like a student learning to use language to
perform complex tasks in high school and college.

In-Context Learning

is like a working professional trying to
figure out his/her manager's instructions

Zero Shot vs Few Shot

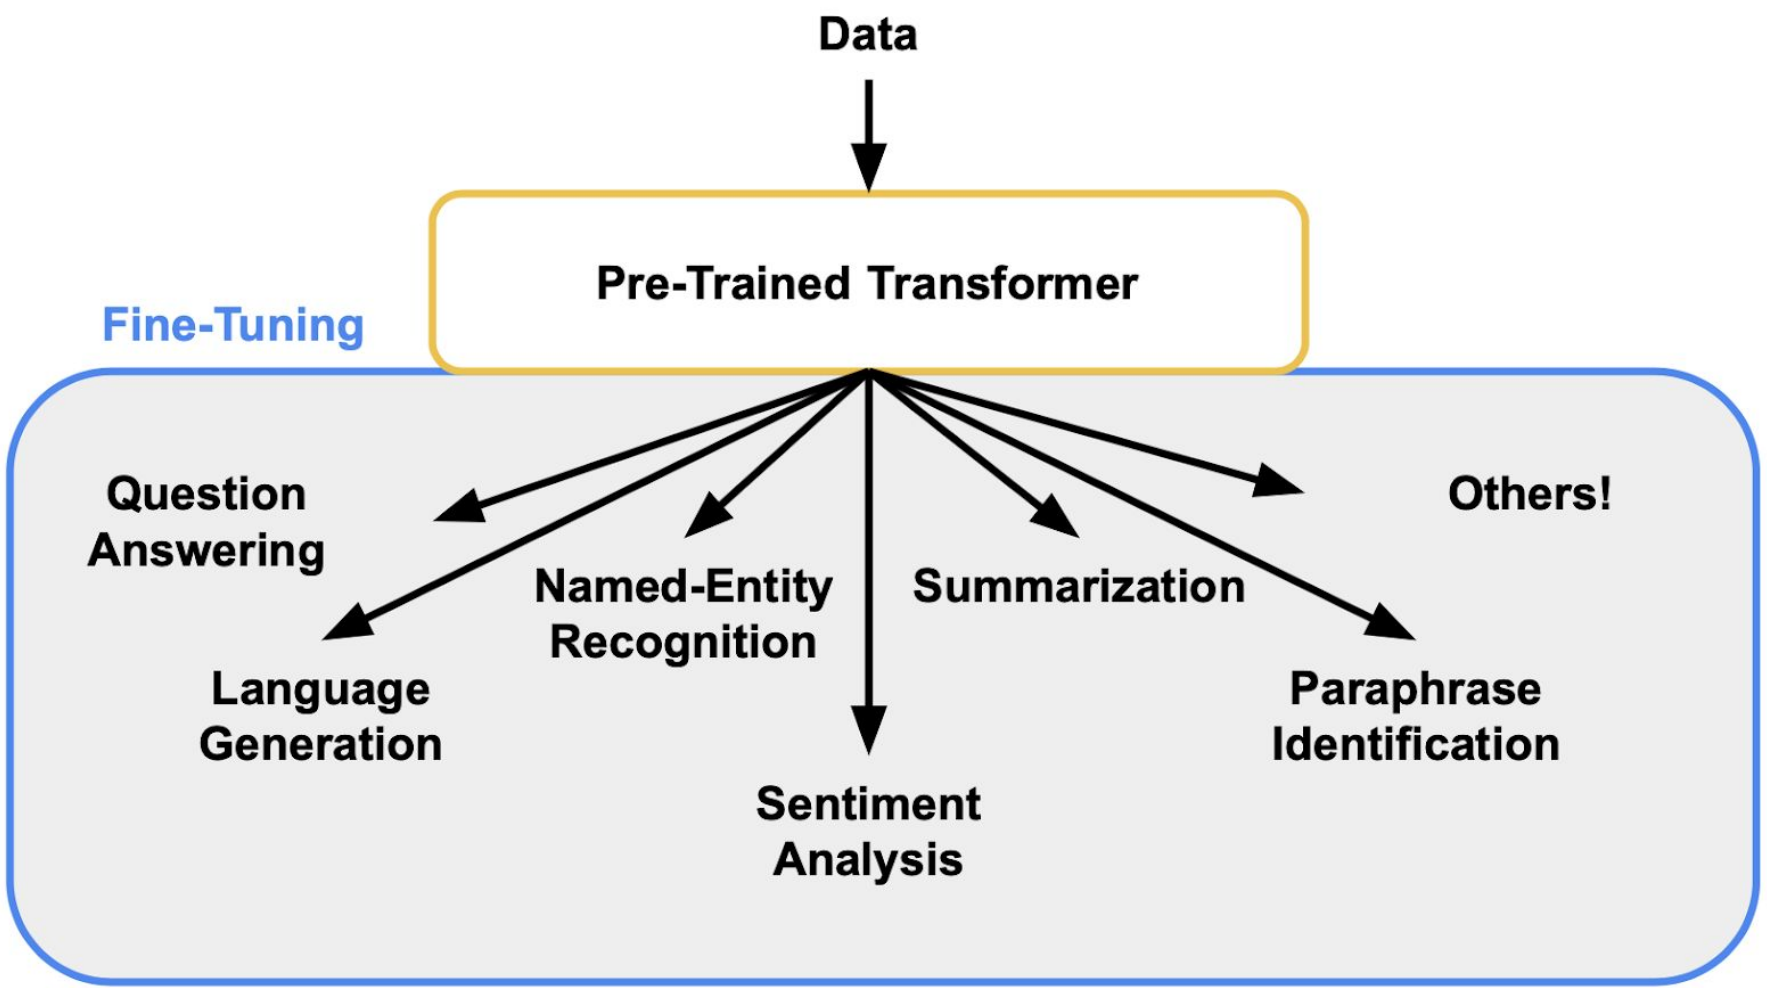
Data
Pre-Trained Transformer
Fine-Tuning
Question Answering
Named-Entity Recognition
Summarization
Others!
Language Generation
Paraphrase Identification
Sentiment Analysis

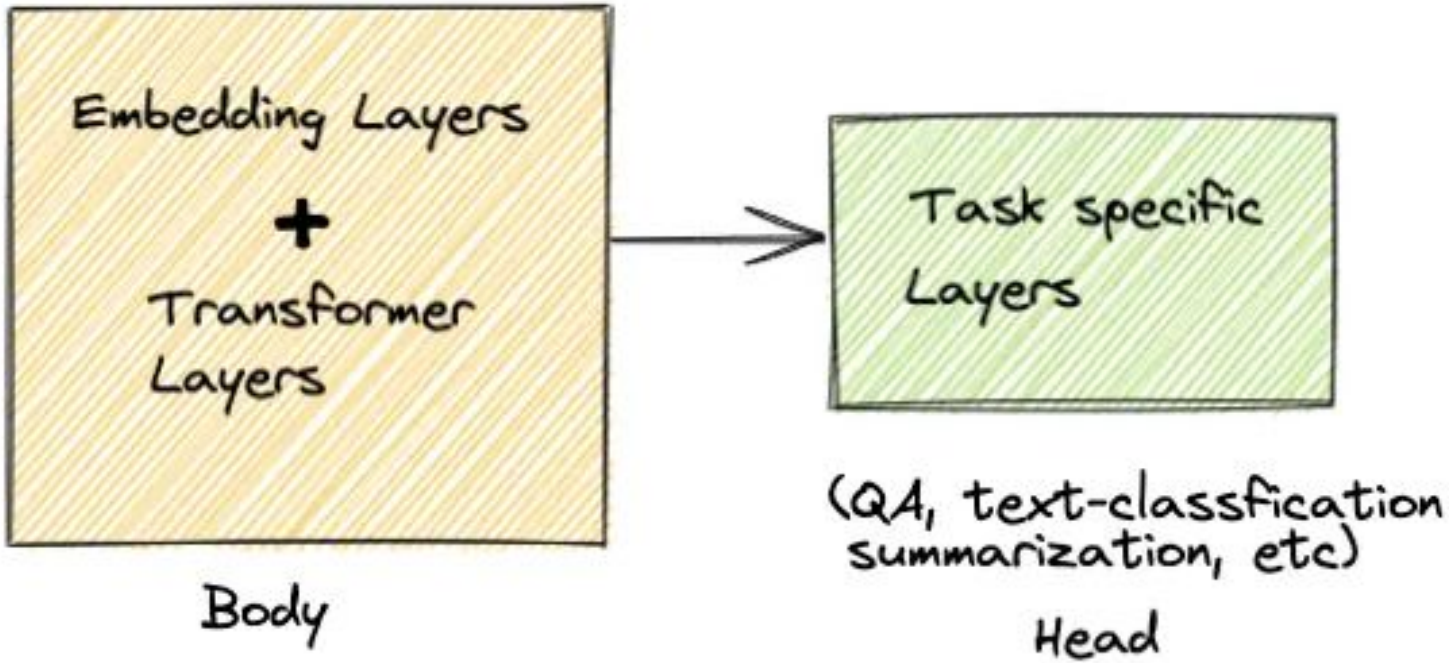
Embedding Layers
+
Transformer Layers
Body
Task specific Layers
(QA, text-classfication summarization, etc)
Head

TEXT CLASSIFICATION

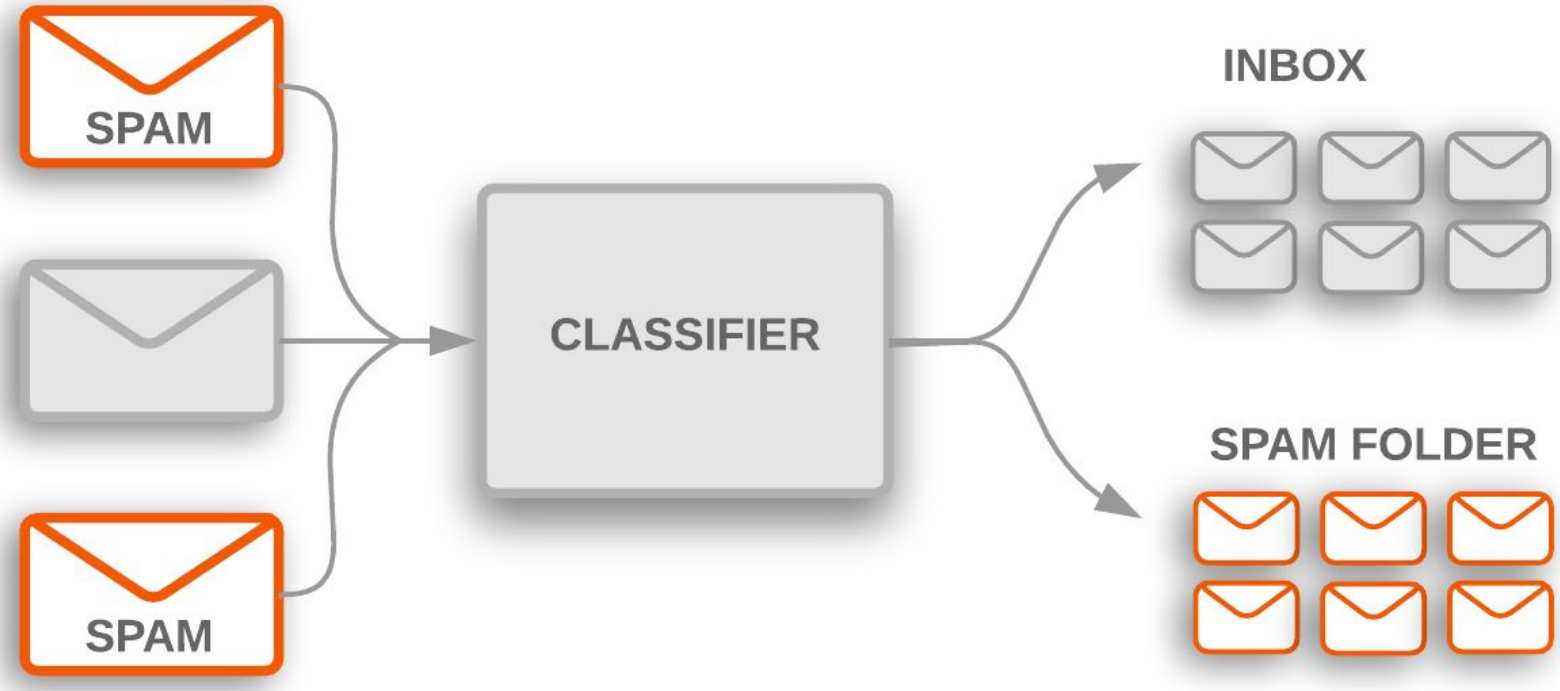
SPAM
SPAM
CLASSIFIER
INBOX
SPAM FOLDER

My experience so far has been fantastic!

POSITIVE

The product is okay I guess.

NEUTRAL

Your support team is useless.

NEGATIVE

Classical NLP Approach

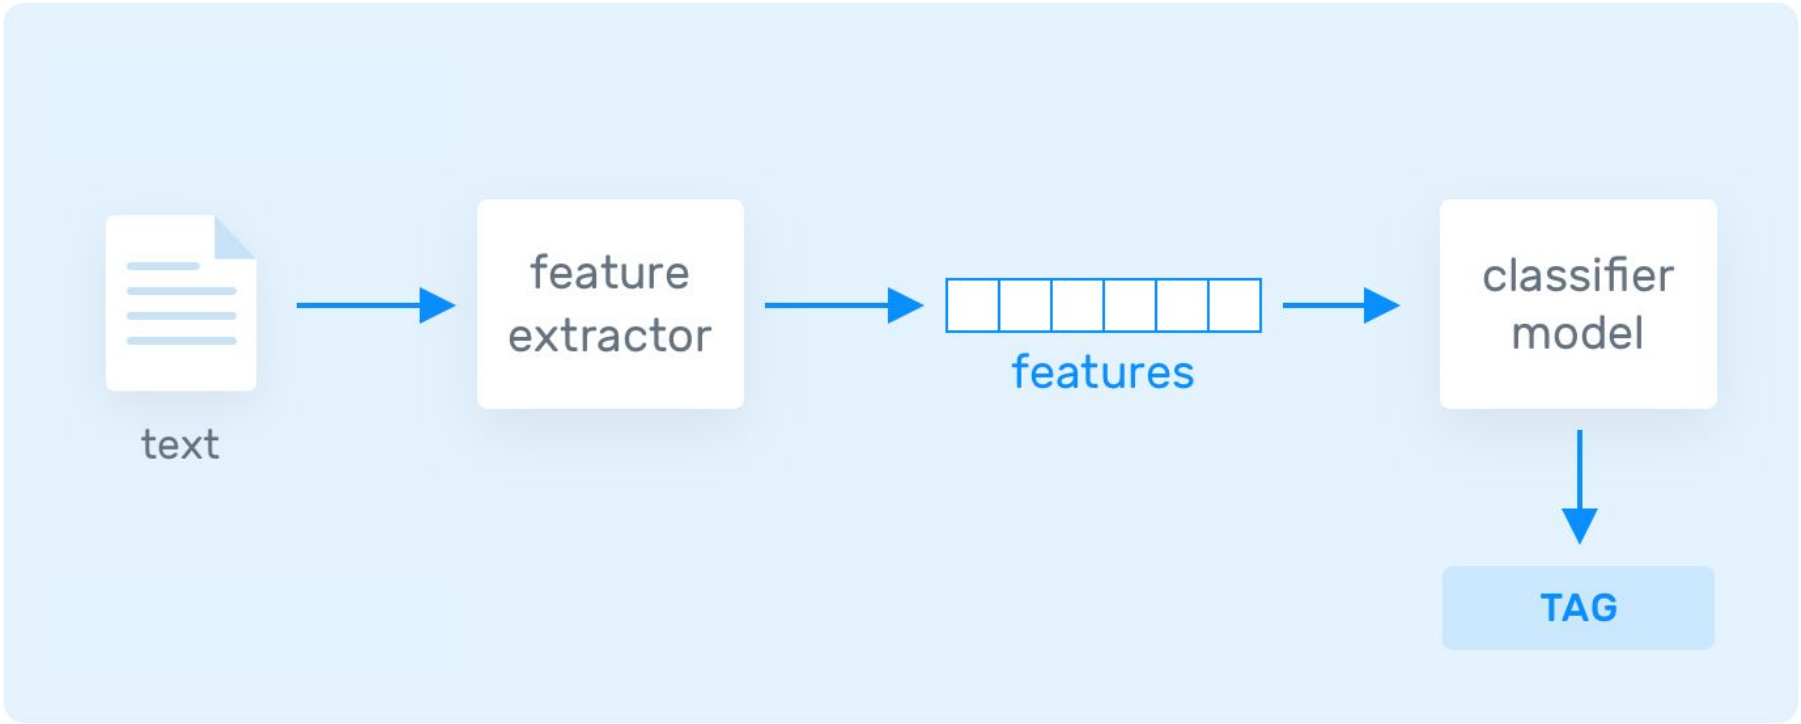

" I like strawberries ", 3 words

↓ ①

"[CLS]", " I ", "like", "straw", "##berries", "[SEP]", 6 tokens

↓ ②

↓ result

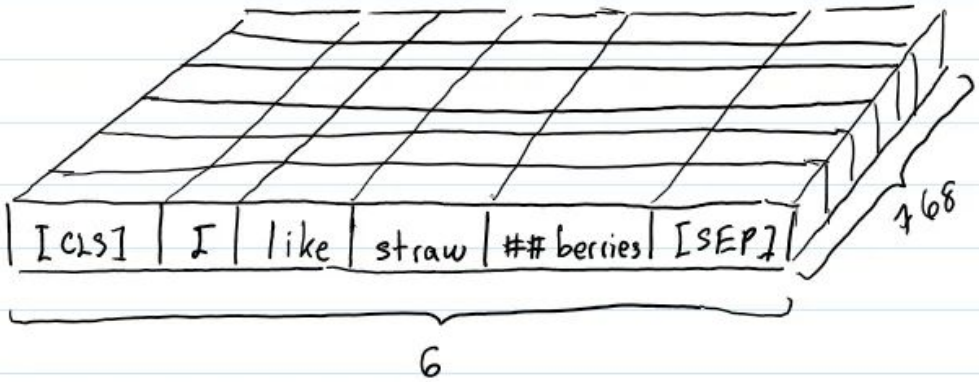

Requires Fine Tuning

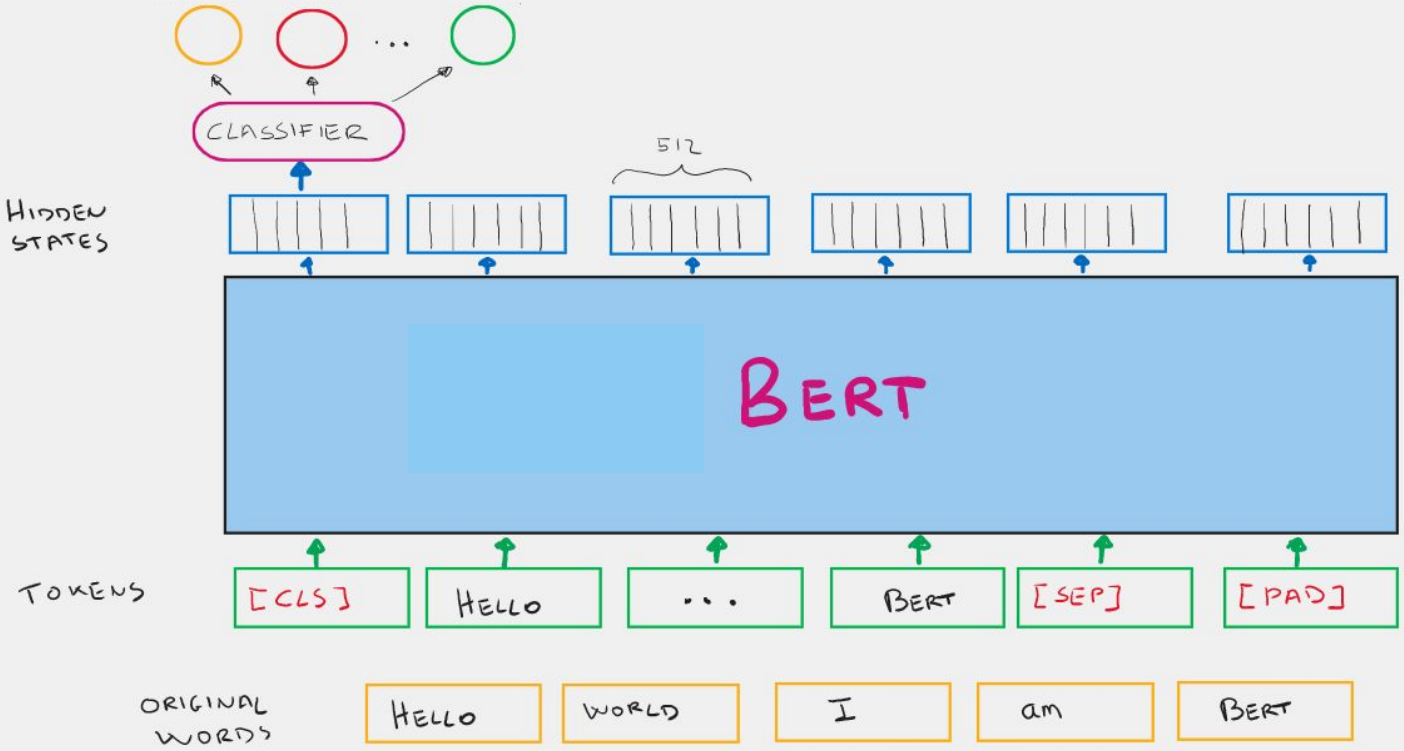

Requires Fine Tuning

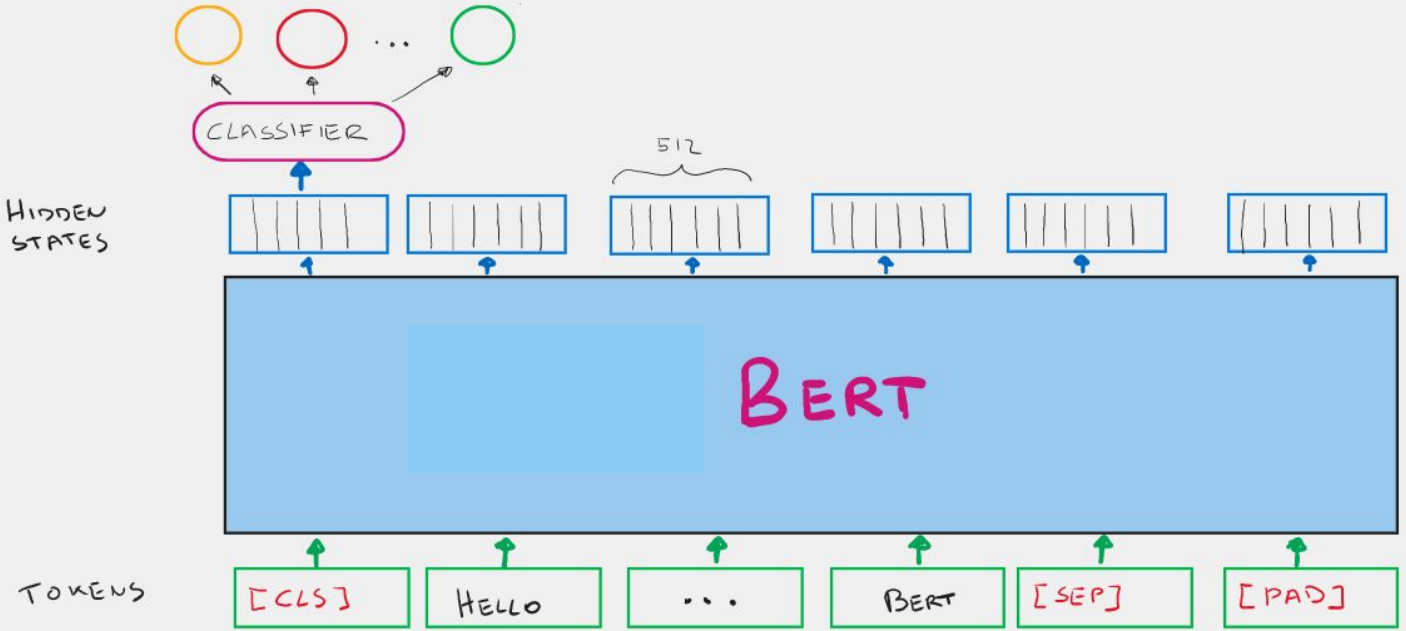

Is only the classifier layer on top trained or
are the BERT parameters also updated during fine-tuning?

NAMED ENTITY RECOGNITION

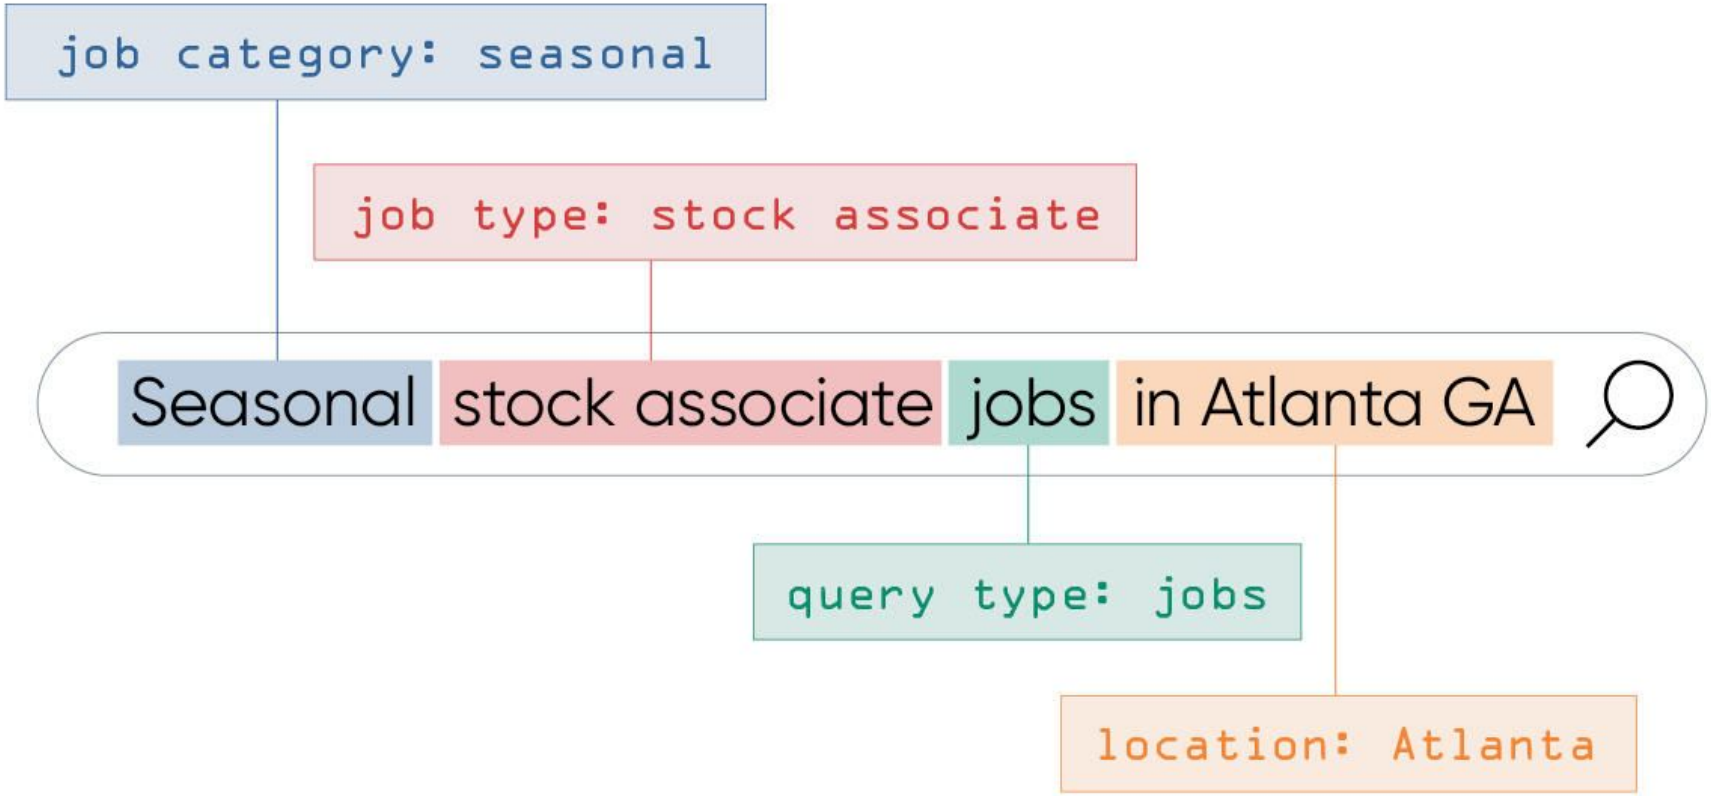
job category: seasonal
job type: stock associate
Seasonal
stock associate
jobs
in Atlanta GA
query type: jobs
location: Atlanta

Ribavirin UMLS: C0035525 was also evaluated against **SARS-CoV-2 infection** , but the **antiviral** UMLS: C0003451

MEDICATION_NAME DIAGNOSIS MEDICATION_CLASS

property of **drugs** UMLS: C0013227 is still not well established against the **SARS-CoV-2** UMLS: C5203670 **negation** .

TREATMENT_NAME DIAGNOSIS

In addition, after **oral** administration, the drug was rapidly absorbed into the **GI tract** UMLS: C0017189 .

ROUTE_OR_MODE BODY_STRUCTURE

The drug has **oral bioavailability** around **64** **%** with large volume of distribution.

ROUTE_OR_MODE EXAMINATION_VALUE EXAMINATION_UNIT

BERT NER : The B-I-O Notation

Yesterday	,	Rohan	Sharma	traveled	to	Mumbai	.
O	O	B-PER	I-PER	O	O	B-LOC	O

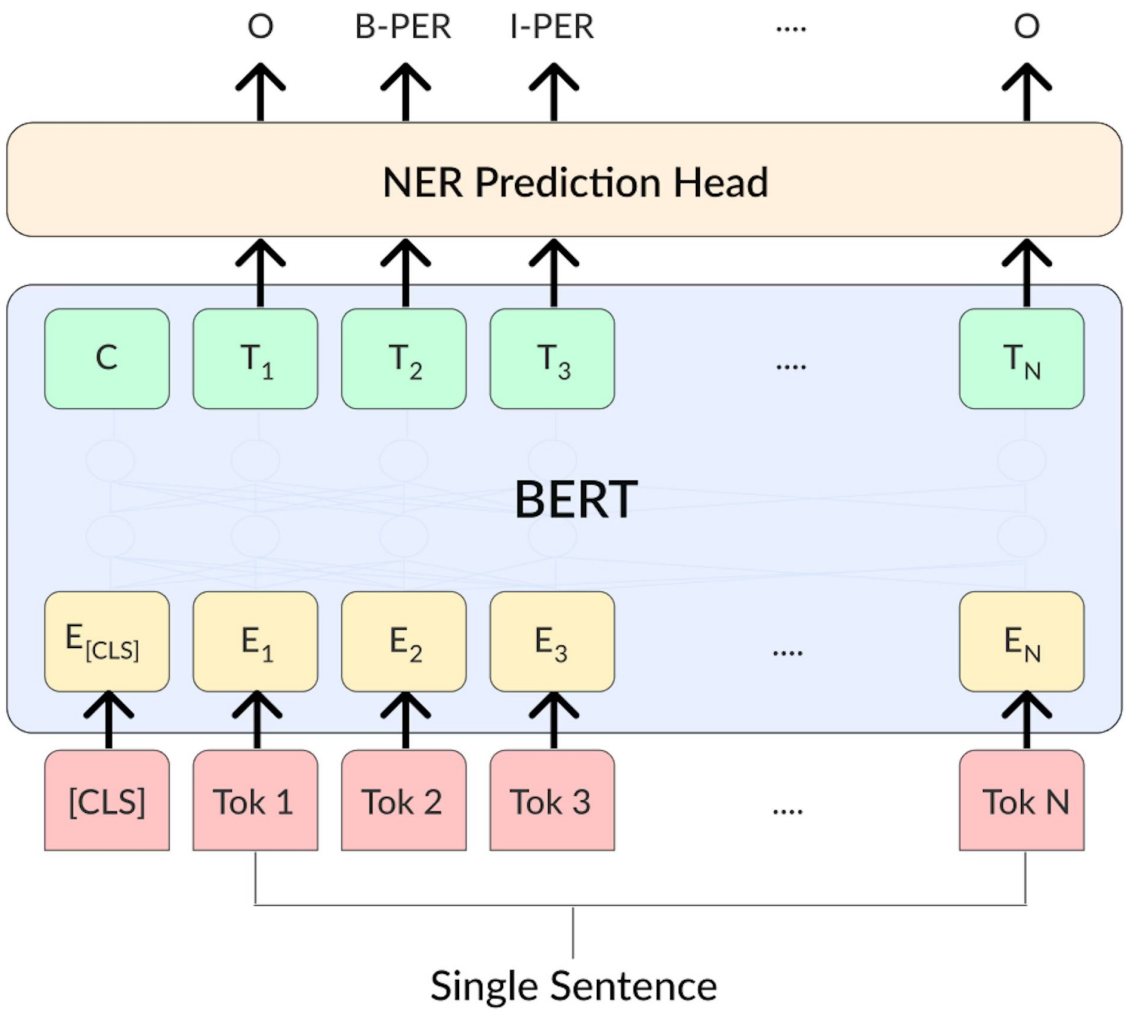
O
B-PER
I-PER
....
O
NER Prediction Head
C
T1
T2
T3
....
TN
BERT
E[CLS]
E1
E2
E3
....
EN
[CLS]
Tok 1
Tok 2
Tok 3
....
Tok N
Single Sentence

INFORMATION RETRIEVAL

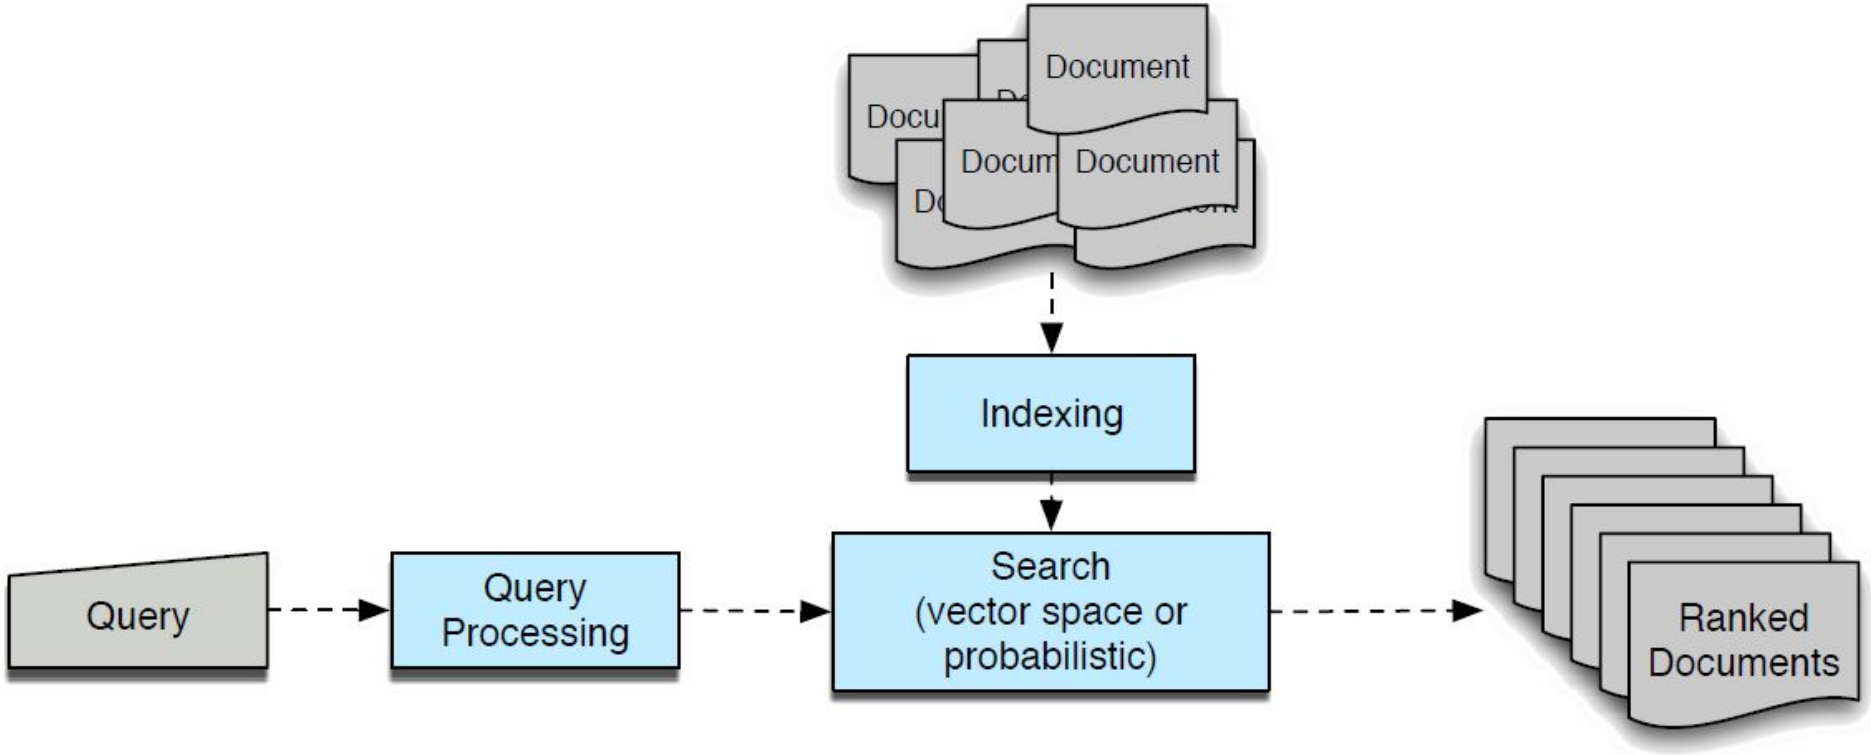
Document
Document
Indexing
Query
Query Processing
Search (vector space or probabilistic)
Ranked Documents

SBERT Fine-Tuning

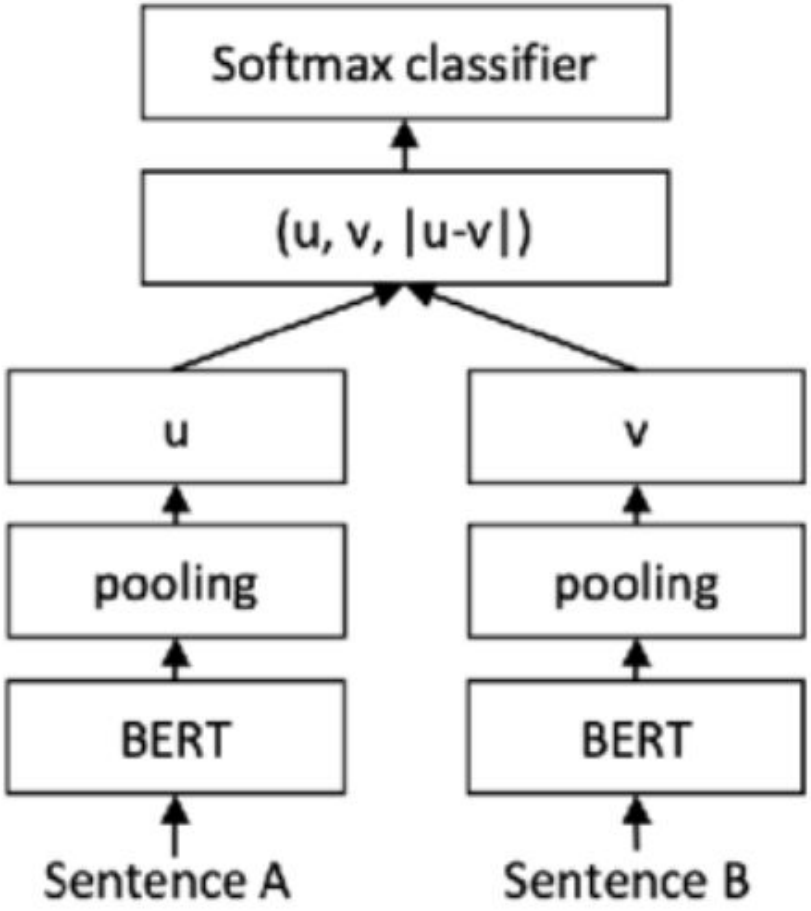

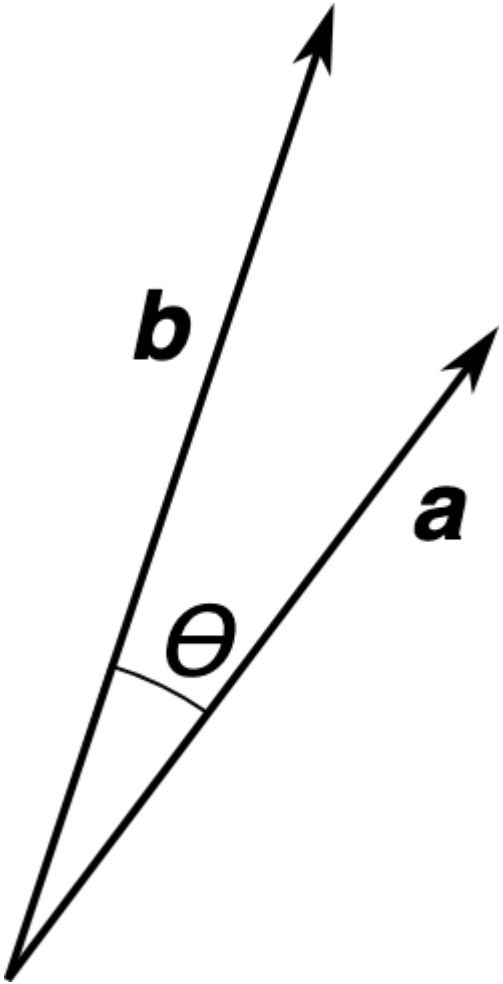

- The query has a vector representation using embeddings
- Documents in the database stored as embeddings

- **Brute Force Approach:**
  Do a dot product of the query vector with the embeddings of all the documents, and choose the one that gives the closest match

- **Hierarchical Navigable Small World (HNSW):**
  Create a layered graph structure of the document embedding vectors so that the search process is made much faster

QUESTION ANSWERING

Passage Sentence

In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under gravity.

Question

What causes precipitation to fall?

Answer Candidate

Gravity

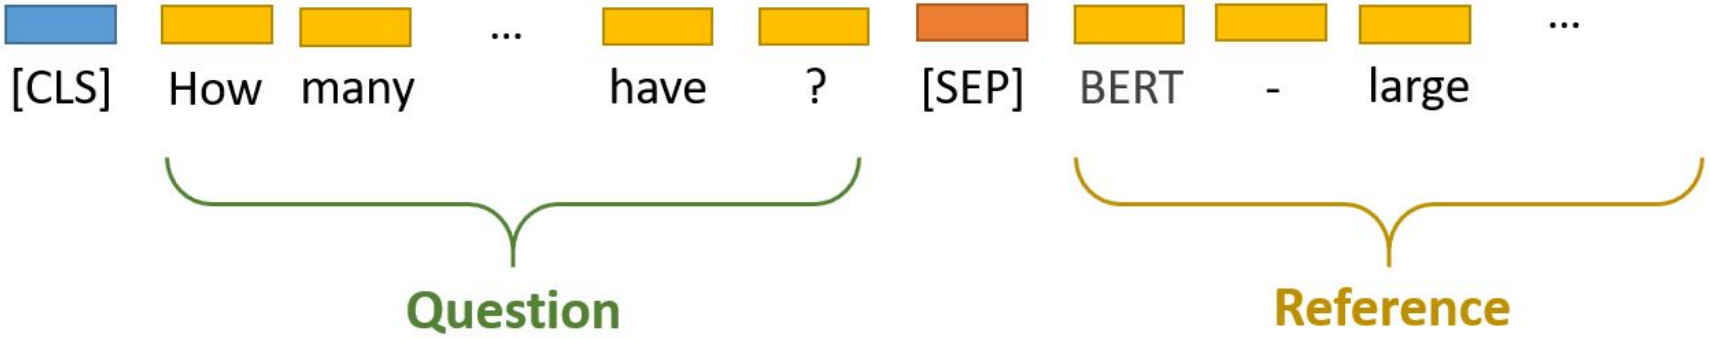

Question: How many parameters does BERT-large have?

Reference Text: BERT-large is really big... it has 24 layers and an embedding size of 1,024, for a total of 340M parameters! Altogether it is 1.34GB, so expect it to take a couple minutes to download to your Colab instance.

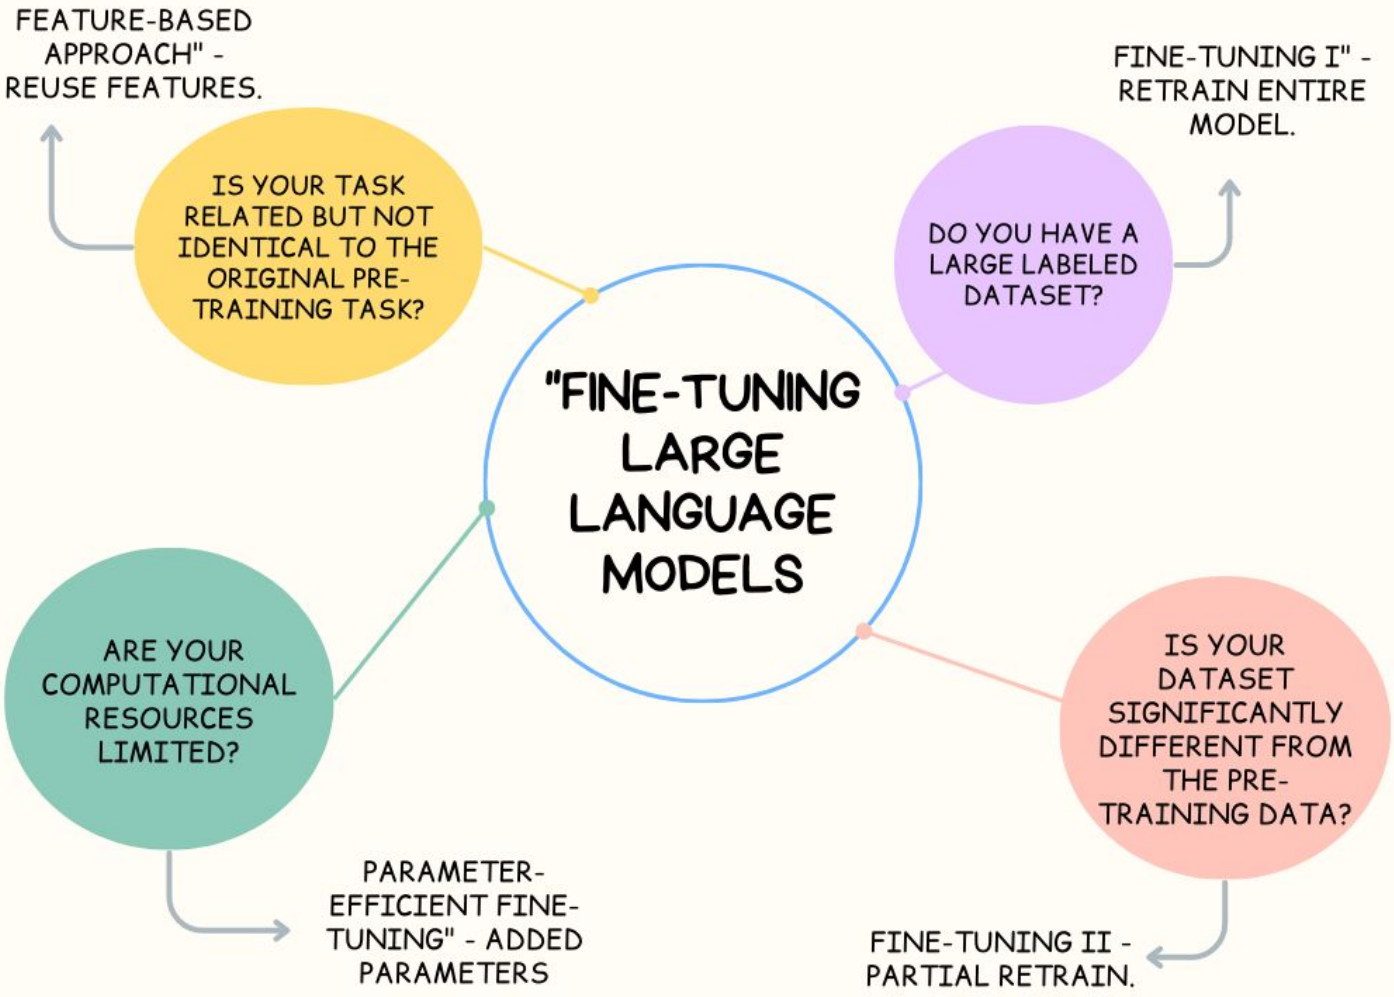
"FINE-TUNING LARGE LANGUAGE MODELS
IS YOUR TASK RELATED BUT NOT IDENTICAL TO THE ORIGINAL PRE-TRAINING TASK?
FEATURE-BASED APPROACH" - REUSE FEATURES.
DO YOU HAVE A LARGE LABELED DATASET?
FINE-TUNING I" - RETRAIN ENTIRE MODEL.
ARE YOUR COMPUTATIONAL RESOURCES LIMITED?
PARAMETER-EFFICIENT FINE-TUNING" - ADDED PARAMETERS
IS YOUR DATASET SIGNIFICANTLY DIFFERENT FROM THE PRE-TRAINING DATA?
FINE-TUNING II - PARTIAL RETRAIN.

How to fine-tune BIG models?

Quantization

- LLMs require a large amount of expensive GPU memory
  - Large number of parameters
  - High precision of the floating point numbers

Model	Original Size	Quantized Size (4-bit)
LLaMA2 7B	13 GB	3.9 GB
LLaMA2 13B	24 GB	7.8 GB
LLaMA2 30B	60 GB	19.5 GB
LLaMA2 65B	120 GB	38.5 GB

NVIDIA A100 has 80 GB memory and costs around INR 12-15 lakhs

Distillation

- Transfer of knowledge from larger “teacher” model
  to a smaller “student” model

- Smaller model represents the bigger model for specific tasks

- Larger model learns the distribution from the data

- Smaller model learns the distribution from the larger model

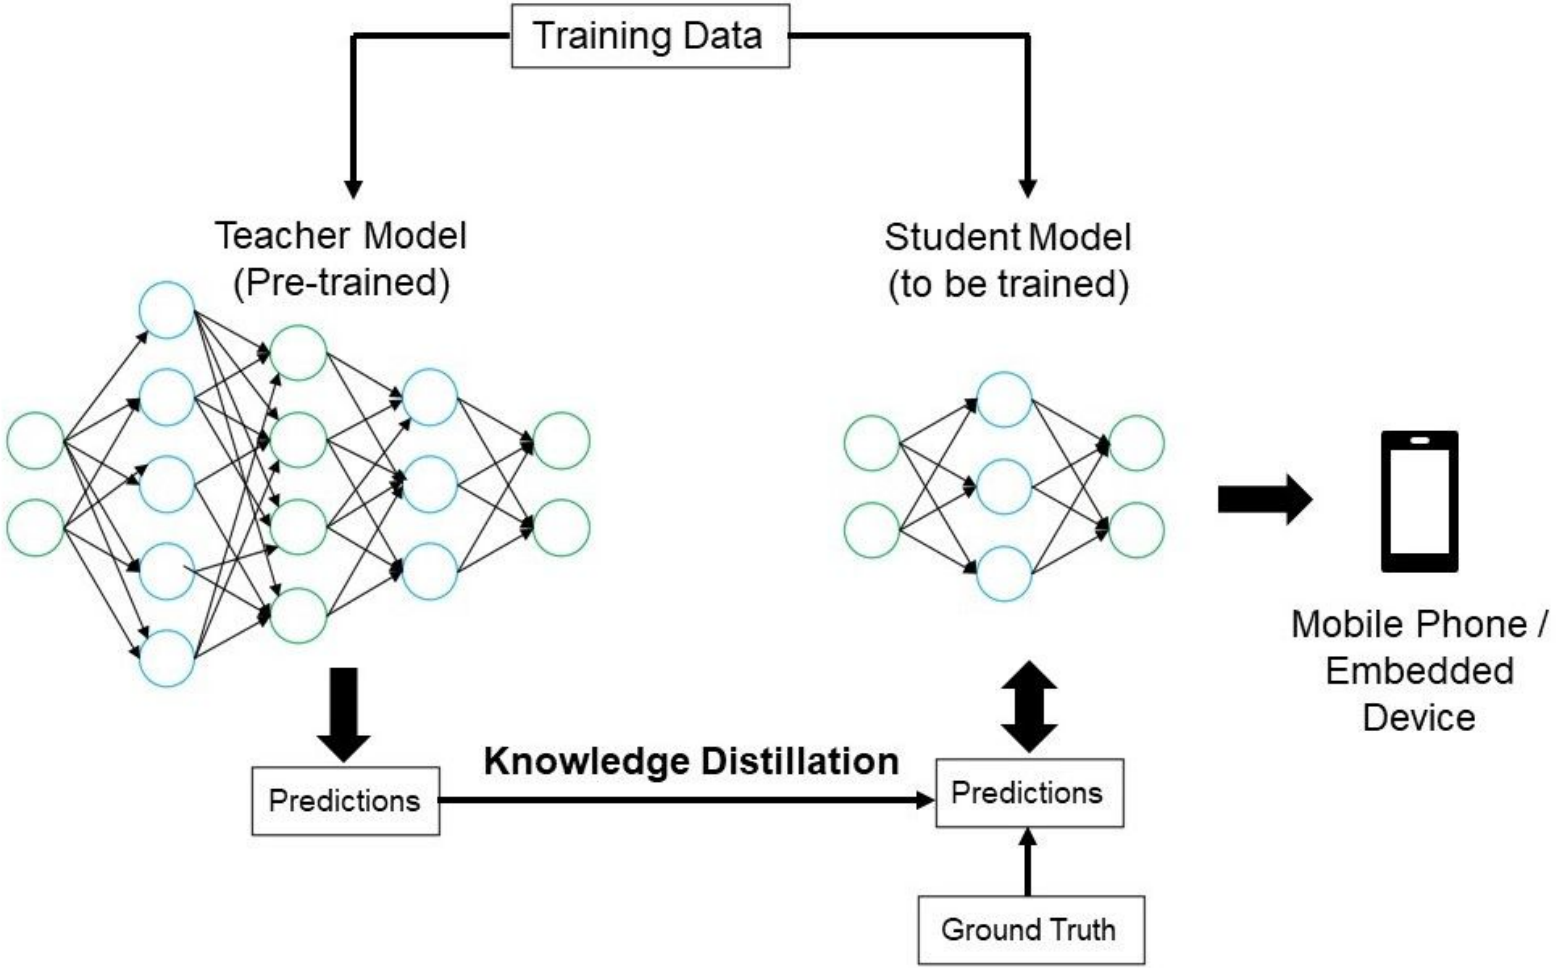
Training Data
Teacher Model
(Pre-trained)
Student Model
(to be trained)
Mobile Phone /
Embedded
Device
Predictions
Knowledge Distillation
Predictions
Ground Truth

	BERT	RoBERT	DistilBERT	XLNet
Size (millions)	**Base**: 110 **Large**: 340	**Base**: 110 **Large**: 340	**Base:** 66	**Base**: ~110 **Large**: ~340
Training Time	**Base**: 8 x V100 x 12 days* **Large:** 64 TPU Chips x 4 days (or 280 x V100 x 1 days*)	**Large:** 1024 x V100 x 1 day; 4-5 times more than BERT.	**Base**: 8 x V100 x 3.5 days; 4 times less than BERT.	**Large:** 512 TPU Chips x 2.5 days; 5 times more than BERT.
Performance	Outperforms state-of-the-art in Oct 2018	2-20% improvement over BERT	5% degradation from BERT	2-15% improvement over BERT
Data	16 GB BERT data (Books Corpus + Wikipedia). 3.3 Billion words.	160 GB (16 GB BERT data + 144 GB additional)	16 GB BERT data. 3.3 Billion words.	**Base**: 16 GB BERT data **Large**: 113 GB (16 GB BERT data + 97 GB additional). 33 Billion words.
Method	BERT (Bidirectional Transformer with MLM and NSP)	BERT without NSP**	BERT Distillation	Bidirectional Transformer with Permutation based modeling

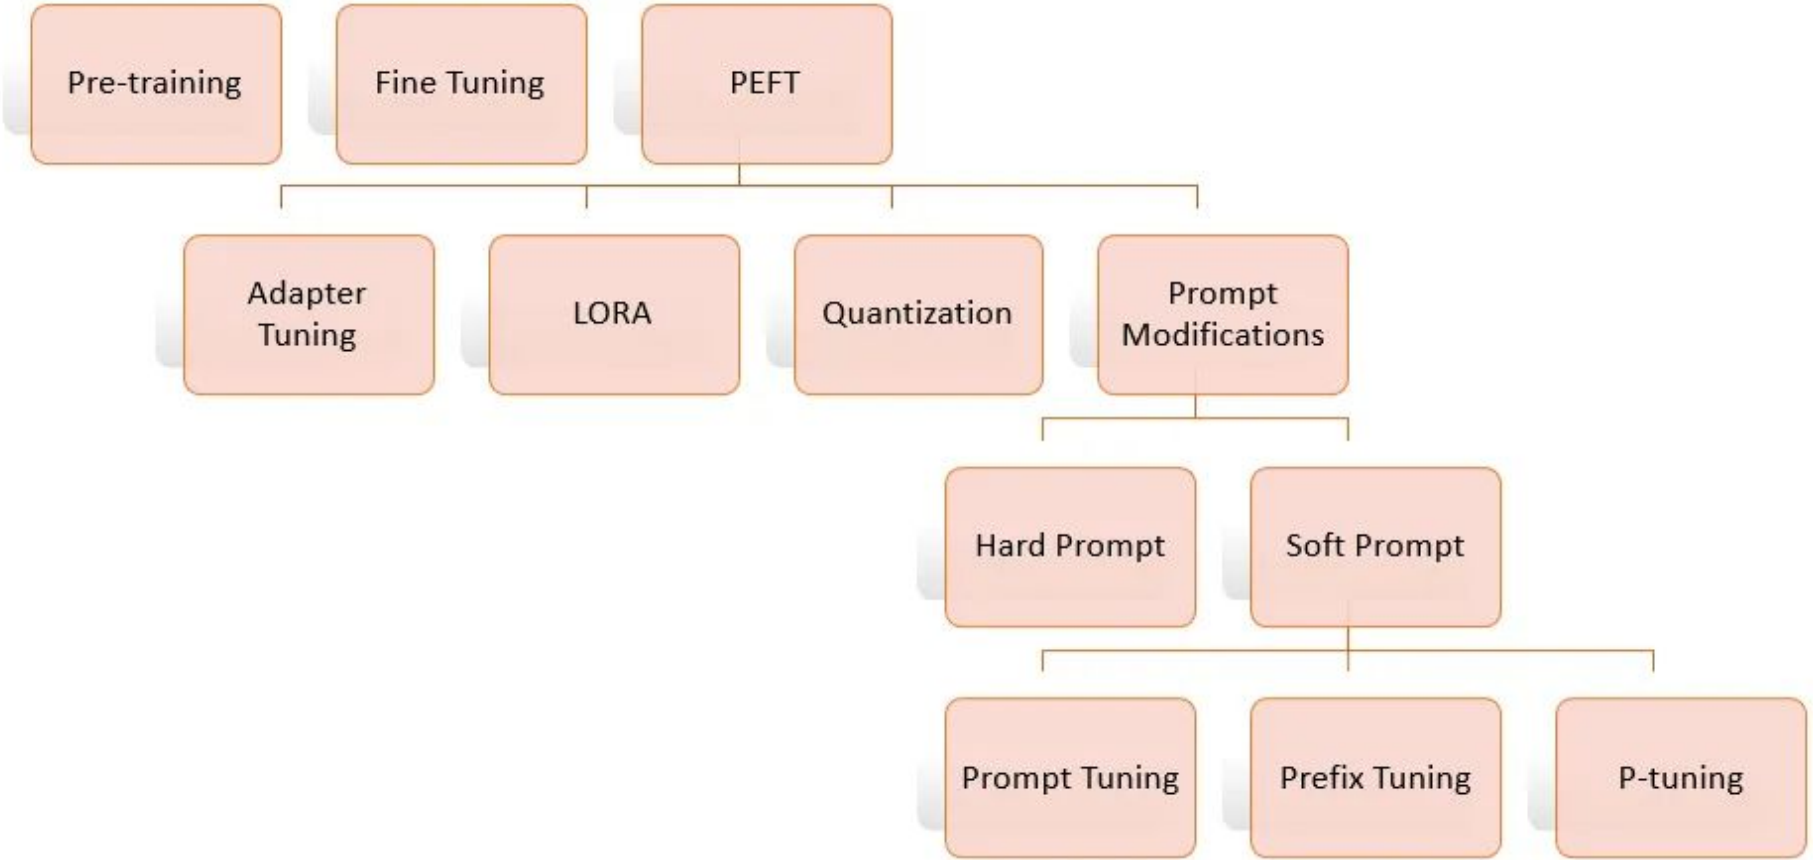
Pre-training
Fine Tuning
PEFT
Adapter Tuning
LORA
Quantization
Prompt Modifications
Hard Prompt
Soft Prompt
Prompt Tuning
Prefix Tuning
P-tuning